

# CSI 695: Scientific Databases

Fall Term 2017

## Lecture 5: Introduction to Searching in Scientific databases


Lectures: Prof. Dr. Matthias Renz

Exercises: TBA

## Recap: Searching in Scientific Databases: Introduction

- Managing **scientific data** usually requires methods that go beyond the capabilities of standard database management systems.
- Scientific data often involves **complex structured data** not well supported by the table schema used in standard databases.
- **Exact match queries** as provided in standard database management systems often **do not suffice for searching in scientific data**, we need something different!!!

Example:  
How would you store/manage molecules in a standard relational database ?




Think about it ...

# Recap: Searching in Scientific Databases: Introduction

- Scientific data often consists of complex structured data including spatial, temporal, spatio-temporal, and multi-media data?

- Spatial data: 1D, 2D, 3D




- Temporal data: time series



- Spatio-temporal data:

- objects moving in a given space



- Multi-media data:

- audio sequences, video sequences



# Recap: Searching in Scientific Databases: Introduction

## ■ A motivating example

- Given an archive with 2,000,000 images (2D objects)
- Is a given image included in that archive?



## □ Challenges

- “included” does not necessarily mean “identical binary representation!”
- Images may vary in
  - Size (scaling, resolution)
  - Perspective (reflection, ...)
  - Coloring, shading
  - Clipping, cutting
  - Add-ons (border, annotation, ...)

## Recap: Searching in Scientific Databases: Introduction

---

- Instead of searching for

**exact matches** (supported by standard databases (relational-, object-relational DBMS))

we need to search for

**“similar” objects!** → new concepts for managing data required (non-standard databases)

# Recap: Searching in Scientific Databases: Introduction

---

- Searching:
  - General problems when searching for similar objects
    - Informal level
      - Similarity depends on the application
        - Searching for images showing “sunset” => **color** is important
        - Searching for images showing “animals” => **shape** is important
      - Similarity depends on the user’s notion
    - Formal level
      - How are the objects **represented**?
      - How can the **similarity** between objects be **modeled**?
    - Pragmatic level
      - Efficient algorithm for **computing the similarity**
      - Efficient algorithm for **searching** in a large disc-resident database


# Recap: Searching in Scientific Databases: Introduction

---

- Searching:
  - Here, we focus on the sub-problem
    - Efficiently **searching for similar objects** in a large database and to some extent on
    - Efficiently **computing the similarity** between objects
  - We assume a very common model of similarity: ***Feature-based similarity***
- Feature-based similarity
  - How can we model the similarity between complex objects like images, 3D objects, video sequences, etc.?
  - Considerations
    - **Efficiency**: Model should allow efficient query processing => use of index structures should be possible
    - **Generality**: Avoid the necessity to develop algorithms and index structures for each application separately but develop a general way to model similarity
      - We have indexes for
      - Spatial data and multi-dimensional vectors
      - General metric data (objects with a metric distance function)

## Recap: Searching in Scientific Databases: Introduction


### ■ Feature-based similarity (cont.)



Application Domains (Object Space)

#### Feature Transformation

Histogramms  
Shape descriptors  
Fourier coefficients  
etc.



Feature (Vector) Space

- Objects from real world are transformed into multi-dimensional feature vector points
  1. Identify a set (sequence) of (numerical) features from objects that best describe the objects
  2. Build a multidimensional point vector from the set (sequence) of features
  3. Manage the point vectors where each point vector has a link to the detailed object descriptions
  4. (Object) point vectors are efficiently organized (managed) using **appropriate index structures**



## Outline

---

- Searching in Scientific Databases: Introduction
- Feature spaces and proximity measures
- Feature transformation for text data
- Algorithmic Paradigms for Similarity Query Processing

# Feature spaces and proximity measures

---


## Feature space

- Intuitively: a domain with a distance function
- Formally: feature space  $\mathbf{F} = (Dom, dist)$ :
  - $Dom$  is a set of attributes / features
  - $dist$ : a numerical measure of the degree to which the two compared objects differ
    - $dist : Dom \times Dom \rightarrow \mathbb{R}_0^+$
- For all  $x, y$  in  $Dom$ ,  $x \neq y$ ,  $dist$  is required to satisfy the following properties:
  - $dist(x, y) > 0$  (non-negativity)
  - $dist(x, x) = 0$  (reflexivity)

# Feature spaces and proximity measures

## Metric space

- Formally: Metric space  $M = \{Dom, dist\}$ :
  - $M$  is a feature space
    - i.e.,  $dist(x,y) > 0$  (non-negativity) and,
    - $dist(x,x) = 0$  (reflexivity)
  - $dist(x, y) = 0 \Rightarrow x = y$  (strictness)
  - $\forall x, y \in Dom: dist(x, y) = dist(y, x)$  (symmetry)
  - $\forall x, y, z \in Dom : dist(x, z) \leq dist(x, y) + dist(y, z)$  (triangle inequality)
- Measures that satisfy all the above properties are called metrics.



## Feature spaces and proximity measures

### ■ Famous example: Euclidean vector space $E=(Dom, dist)$

□  $(Dom, dist)$  is a metric space

□  $Dom = \mathbb{R}^d$

□  $dist(x, y) = \sqrt{\sum_{i=1}^d (x_i - y_i)^2}$

### ■ Notation:

□ Euclidean vector space =: “Feature space”

□ Vectors (Objects in the Euclidean feature space) =: “Feature vectors”

□ The  $d$  dimensions of the vector space =: “Features”


### ■ Standardization is necessary, if scales differ!

□ e.g., age (e.g., range [0-100] vs salary (e.g., range: 10000-100000))

*We will come back to this in a few slides*

point	x	y
p1	0	2
p2	2	0
p3	3	1
p4	5	1

Point coordinates



	p1	p2	p3	p4
p1	0	2.828	3.162	5.099
p2	2.828	0	1.414	3.162
p3	3.162	1.414	0	2
p4	5.099	3.162	2	0

Distance matrix

## Feature spaces and proximity measures

---


- Manhattan distance or City-block distance ( $L_1$  norm)
  - $dist_1 = |p_1 - q_1| + |p_2 - q_2| + \dots + |p_d - q_d|$
  - The sum of the absolute differences of the  $p, q$  coordinates
- Euclidean distance ( $L_2$  norm)
  - $dist_2 = ((p_1 - q_1)^2 + (p_2 - q_2)^2 + \dots + (p_d - q_d)^2)^{1/2}$
  - The length of the line segment connecting  $p$  and  $q$
- Supremum distance ( $L_{max}$  norm or  $L_\infty$  norm)
  - $dist_\infty = \max\{|p_1 - q_1|, |p_2 - q_2|, \dots, |p_d - q_d|\}$
  - The max difference between any attributes of the objects.
- Minkowski Distance (Generalization of  $L_p$ -distance)
  - $dist_p = (|p_1 - q_1|^p + |p_2 - q_2|^p + \dots + |p_d - q_d|^p)^{1/p}$

## Feature spaces and proximity measures

### ■ Example

point	x	y
p1	0	2
p2	2	0
p3	3	1
p4	5	1

Point coordinates



L1	p1	p2	p3	p4
p1	0	4	4	6
p2	4	0	2	4
p3	4	2	0	2
p4	6	4	2	0

$L_1$  distance matrix


L2	p1	p2	p3	p4
p1	0	2.828	3.162	5.099
p2	2.828	0	1.414	3.162
p3	3.162	1.414	0	2
p4	5.099	3.162	2	0

$L_2$  distance matrix

$L_\infty$	p1	p2	p3	p4
p1	0	2	3	5
p2	2	0	1	3
p3	3	1	0	2
p4	5	3	2	0

$L_\infty$  distance matrix


## Feature spaces and proximity measures



Source: <http://www.econ.upf.edu/~michael/stanford/maeb5.pdf>

## Feature spaces and proximity measures

- Let  $x, y$  in  $[-1, 1]$
- For L1 norm
  - $|(x, y)|_1 = 1 \Rightarrow x + y = 1$
  - If  $x = 1, y = 0$
  - If  $x = 0.8, y = 0.2$
  - ...
- For L2 norm
  - $(x^2 + y^2)^{1/2} = 1$
  - It is circle
- ...



Unit Circle for different Lp-distances

Source: <https://de.wikipedia.org/wiki/P-Norm>



# Normalization

- Attributes with large ranges outweigh ones with small ranges
  - e.g. income [10K-100K]; age [10-100]
- To balance the “contribution” of an attribute  $A$  in the resulting distance, the attributes are scaled to fall within a small, specified range.
- min-max normalization: to  $[new\_min_A, new\_max_A]$

$$v' = \frac{v - min_A}{max_A - min_A} (new\_max_A - new\_min_A) + new\_min_A$$

- e.g. normalize age=30 in [0-1], when min=10,max=100.  $new\_age = ((30-10)/(100-10)) * (1-0) + 0 = 2/9$
- z-score normalization also called zero-mean normalization
  - After zero-mean normalizing each feature will have a mean value of 0

$$v' = \frac{v - mean_A}{stand\_dev_A}$$

e.g. normalize 70,000 iff  $\mu=50,000$ ,  $\sigma=15,000$ .  
 $new\_value = (70,000 - 50,000) / 15,000 = 1.33$

## Proximity between binary attributes 1/2

- A binary attribute has only two states: 0 (absence), 1 (presence)
- A contingency table for binary data

		<i>Instance j</i>		
		1	0	sum
<i>Instance i</i>	1	$q$	$r$	$q + r$
	0	$s$	$t$	$s + t$
	sum	$q + s$	$r + t$	$p$

$q$  = the number of attributes where  $i$  was 1 and  $j$  was 1  
 $t$  = the number of attributes where  $i$  was 0 and  $j$  was 0

$s$  = the number of attributes where  $i$  was 0 and  $j$  was 1  
 $r$  = the number of attributes where  $i$  was 1 and  $j$  was 0

- Simple matching coefficient  
(for symmetric binary variables)

$$d(i, j) = \frac{r + s}{q + r + s + t}$$

- for asymmetric binary variables:

$$d(i, j) = \frac{r + s}{q + r + s}$$

- Jaccard coefficient  
(for *asymmetric* binary variables)

$$\text{sim}_{\text{Jaccard}}(i, j) = \frac{q}{q + r + s}$$

## Proximity between binary attributes 2/2

■ Example:

Name	Fever	Cough	Test-1	Test-2	Test-3	Test-4
Jack	1	0	1	0	0	0
Mary	1	0	1	0	1	0
Jim	1	1	0	0	0	0

$$d(jack, mary) = \frac{0 + 1}{2 + 0 + 1} = 0.33$$

$$d(jack, jim) = \frac{1 + 1}{1 + 1 + 1} = 0.67$$

$$d(jim, mary) = \frac{1 + 2}{1 + 1 + 2} = 0.75$$

(from previous slide)

q = the number of attributes where i was 1 and j was 1  
t = the number of attributes where i was 0 and j was 0

s = the number of attributes where i was 0 and j was 1  
r = the number of attributes where i was 1 and j was 0

$$d(i, j) = \frac{r + s}{q + r + s}$$

## Proximity between categorical attributes

- A nominal attribute has >2 states (generalization of a binary attribute)

- e.g. color={red, blue, green}

- Method 1: Simple matching

- m: # of matches, p: total # of variables

$$d(i, j) = \frac{p - m}{p}$$

Name	Hair color	Occupation
Jack	Brown	Student
Mary	Blond	Student
Jim	Brown	Architect

- Method 2: Map it to binary variables

- create a new binary attribute for each of the  $M$  nominal states of the attribute

Name	Brown hair	Blond hair	IsStudent	IsArchitect
Jack	1	0	1	0
Mary	0	1	1	0
Jim	1	0	0	1

## Selecting the right proximity measure

---

- The proximity function should fit the type of data
  - For dense continuous data, metric distance functions like Euclidean are often used.
  - For sparse data, typically measures that ignore 0-0 matches are employed
    - We care about characteristics that objects share, not about those that both lack
- Domain expertise is important, maybe there is already a state-of-the-art proximity function in a specific domain and we don't need to answer that question again.
- In general, choosing the right proximity measure can be a very time consuming task
- Other important aspects: How to combine proximities for heterogeneous attributes (binary and numeric and nominal etc.)
  - e.g., using attribute weights

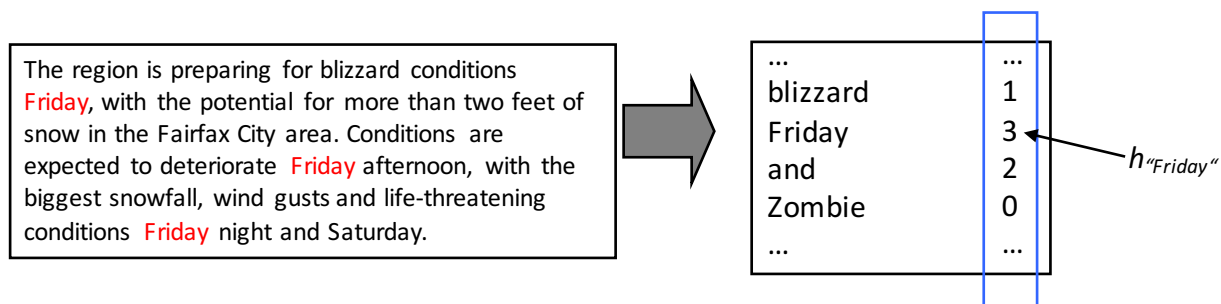
## Outline

---

- Searching in Scientific Databases: Introduction
- Feature spaces and proximity measures
- Feature transformation for text data
- Algorithmic Paradigms for Similarity Query Processing

## Feature transformations for text data 1/6

- Text represented as a set of terms (“Bag-Of-Words” model)
  - Terms:
    - Single words (“cluster“, “analysis“..)  
or
    - bigrams, trigrams, ...n-grams (“cluster analysis“..)
  - Transformation of a document  $d$  in a vector  $r(d) = (h_1, \dots, h_d)$ ,  $h_i \geq 0$ : the frequency of term  $t_i$  in  $d$



## Feature transformations for text data 2/6

---

- Challenges/Problems in Text Mining:
  1. Common words (“e.g.”, “the”, “and”, “for”, “me”)
  2. Words with the same root (“fish”, “fisher”, “fishing”,...)
  3. Very high-dimensional space (dimensionality  $d > 10.000$ )
  4. Not all terms are equally important
  5. Most term frequencies  $h_i = 0$  (“sparse feature space”)
- More challenges due to language:
  - Different words have same meaning (synonyms)
    - “freedom” – “liberty”
  - Words have more than one meanings
    - e.g. “java”, “mouse”



## Feature transformations for text data 3/6

---

- Problem 1: Common words (“e.g.”, “the”, “and”, “for”, “me”)
  - Solution: ignore these terms (Stopwords)

There are stopwords list for all languages in WWW.
  
- Problem 2: Words with the same root (“fish”, “fisher”, “fishing”,...)
  - Solution: Stemming

Map the words to their root

    - “fishing”, “fished”, “fish”, and “fisher” to the root word, “fish”.

For English, the Porter stemmer is widely used.  
(Porters Stemming Algorithms: <http://tartarus.org/~martin/PorterStemmer/index.html>)

Stemming solutions exist for other languages also.

The root of the words is the output of stemming.

## Feature transformations for text data 4/6

- Problem 3: Too many features/ terms
  - Solution: Select the most important features (“Feature Selection”)
  - Example: average document frequency for a term
    - Very frequent items appear in almost all documents
    - Very rare terms appear in only a few documents

Ranking procedure:

1. Compute document frequency for all terms  $t_i$  :
2. Sort terms w.r.t.  $DF(t_i)$  and get  $rank(t_i)$
3. Sort terms by  $score(t_i) = DF(t_i) \cdot rank(t_i)$   
e.g.  $score(t_{23}) = 0.82 \cdot 1 = 0.82$   
 $score(t_{17}) = 0.65 \cdot 2 = 1.3$
4. Select the  $k$  terms with the largest  $score(t_i)$

$$DF(t_i) = \frac{\#Docs\ containing\ t_i}{\#All\ documents}$$

Rank	Term	DF
1.	$t_{23}$	0.82
2.	$t_{17}$	0.65
3.	$t_{14}$	0.52
4.	...	...

## Feature transformations for text data 5/6

### ■ Problem 4: Not all terms are equally important

- Idea: Very frequent terms are less informative than less frequent words. Define such a term weighting schema.
- Solution: TF-IDF (Term Frequency · Inverse Document Frequency)

Consider both the importance of the term in the document and in the whole collection of documents.

$$TF(t, d) = \frac{n(t, d)}{\sum_{w \in d} n(w, d)} \quad \text{The relative frequency of term } t \text{ in } d \text{ [} n(t, d) = \# t \text{ in } d \text{]}$$

$$IDF(t) = \log\left(\frac{|DB|}{|\{d \mid d \in DB \wedge t \in d\}|}\right) \quad \text{Inverse frequency of term } t \text{ in all DB}$$

$$TF \times IDF = TF(t, d) \cdot IDF(t)$$

Feature vector with TF IDF :  $r(d) = (TF(t_1, d) \cdot IDF(t_1), \dots, TF(t_n, d) \cdot IDF(t_n))$

## Feature transformations for text data 6/6

- Problem 5: for most of the terms  $h_i = 0$ 
  - Euclidean distance is not a good idea: it is influenced by vectors lengths
  - Idea: use more appropriate distance measures

**Jaccard Coefficient:** Ignore terms absent in both documents

$$d_{Jaccard}(d_1, d_2) = 1 - \frac{|d_1 \cap d_2|}{|d_1 \cup d_2|} = \frac{|\{t | t \in d_1 \wedge t \in d_2\}|}{|\{t | t \in d_1 \vee t \in d_2\}|}$$

**Cosine Coefficient:** Consider term values (e.g. TFIDF values)

$$d_{\cosinus}(d_1, d_2) = 1 - \frac{\langle d_1, d_2 \rangle}{\|d_1\| \cdot \|d_2\|} = 1 - \frac{\sum_{i=0}^n (d_{1,i} \cdot d_{2,i})}{\sqrt{\sum_{i=0}^n d_{1,i}^2} \cdot \sqrt{\sum_{i=0}^n d_{2,i}^2}}$$